

Pointer Basics

Optimized C++

Ed Keenan

23 January 2019

13.0.6.3.4 Mayan Long Count



Goals

- Tweak your understanding of pointers
 - Hopefully it's a review
- Learn a cool geek tricks
 - Still great party material
 - Even works at lunches
 - Multi-useful material



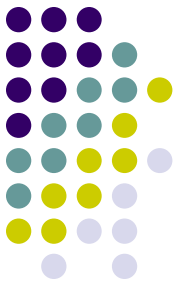
This class is great!



Welcome to the Club!

- Pointers are very important
 - You need to understand completely
 - Be proficient
 - Understanding offsets
 - Understanding arrays
 - And more
- Important Code Ninja Skills
 - Train, Train, Train
 - Drill, Drill, Drill

R Lee Ermey - Wisdom



- Full Metal Jacket
 - Rifleman's Creed



Programmer's Prayer

- This is my compiler.
 - There are many like it, but this one is mine.
 - My compiler is my best friend. It is my life. I must master it, as I must master my life.
- Without me, my compiler is useless.
 - Without my compiler, I am useless.
- I must write my code true.
 - I must debug better than my competition who are trying to take my job.
 - I must release to market before they reach beta.
- I will.



Never under estimate

- Never under estimate
 - A motivated software developer
 - Who knows pointers!
 - They can code anything!



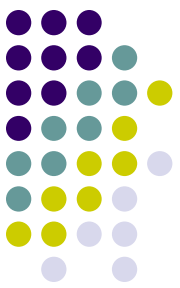
Reality Check!

- I really need to get everyone's attention
 - Yes you in the back of the room
- Here's a Quiz!
 - <http://www.online-stopwatch.com/countdown-clock/full-screen>



Answers

- | | | | |
|-----|------|-----|------------|
| 1. | 0x3F | 12. | 0x3f12cdab |
| 2. | 0xb5 | 13. | 0x11225529 |
| 3. | 0xcd | 14. | 0x35d3b533 |
| 4. | 0x44 | 15. | 0x13a98856 |
| 5. | 0xcd | 16. | 0x11225529 |
| 6. | 0x75 | 17. | 0x0366 |
| 7. | 0x11 | 18. | 0x3375 |
| 8. | 0x29 | 19. | 0x2668 |
| 9. | 0x56 | 20. | 0x8856 |
| 10. | 0xa9 | 21. | 0x13 |
| 11. | 0xa9 | 22. | 0x29 |
| | | 23. | 0x14 |



Pointers == Arrays?

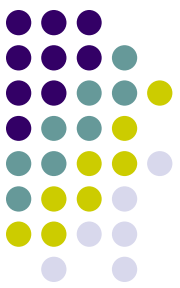
```
int *p;
```

```
p[3] = 5;
```

```
int a[10];
```

```
a[3] = 5;
```

- Are pointer and arrays the same?
 - Both access the 3rd element
 - **p[3]** accesses the 3rd element
 - **a[3]** accesses the 3rd element
 - Syntax looks the same
 - let's look into this a little deeper
 - **a[3]** translates to:
 - ***(a + 3)**
 - **p[3]** translates to:
 - ***(p+3)**



Arrays

```
int a[10];
```


- $*(a + 3)$
 1. dereference the following
 2. address a
 3. plus the $3 * \text{sizeof}(a \text{ type})$
- $*(a + 3 * \text{sizeof}(\text{int}))$
- Array is the literal storage
 - It's the address the first element

a 0x03b1ff80:

a[0]
a[1]
...
...
...
...
a[9]



Pointers



- **Pointer** is a location in memory that contains an address to the memory

```
int *p;
```

- `*(p+3)`
 1. dereference the following
 2. address that is contained in `p`
 3. plus the `3 * sizeof (what p points to)`
- `*(p + 3 * sizeof(int))`



Shocking truth

- Pointers and Arrays are different!
 - Keep those pictures in mind.
- Ask yourself,
 - Is the data or a register that points to the data?



Pointer arithmetic

- $*(p + 3);$
- P is a register that contains an address
- 3 is the *number* of locations of the type
- $3 * \text{sizeof}(\text{what } p \text{ points to})$
- Double $*p;$
 - $*(p + 3)$
 - $*(p + 3 * \text{sizeof}(\text{double}))$
 - $*(p + 3 * 8)$
 - $*(p + 24)$
- Address in variable p + 24 bytes is the address that is dereference



[Bracket]

- $p[3]$
- Translates to
 - $*(p + 3)$
- Addition is Commutative
 - $*(3 + p)$ is same
- Negative numbers are allowed
 - $*(p - 3)$
 - $p[-3]$

- So it party time....
- if
 - $*(p + 3) == p[3]$
- then
 - $*(3 + p) == 3[p]$



How did it work?
Why the K in cool?
Mortal Kombat team.



Casting

- `static_cast<dest type> (source variable)`
 - Converts to the destination
 - Old style cast, `float x = (float) d_int;`
- `reinterpret_cast<dest type> (source variable)`
 - Does NOT convert data.
 - Shut's up compiler warning, just copy data literally
- `dynamic_cast<dest type> (source variable)`
 - Casting down in hierarchy through polymorphism
 - Cast a derived object to a base object
 - Can be dangerous
 - Needs Run-Time Type Information (RTTI)
- `const_cast<dest type> (source variable)`
 - Removes const off of a variable
 - If you use this, generally something is wrong



Pre-increment $*++p$

$*++p$

- Translates to:

1. $p = p + 1;$
2. $*p;$

$*--p$

- Translates to:

1. $p = p - 1;$
2. $*p;$


- Parenthesis does nothing
- $*(++p) == *++p$
- Remember both
 - pre-increment
 - pre-decrement
 - Address is processed **before** dereferencing



Cave Formations

- Stalagmite
 - Forms for deposits starting from the floor
 - Be careful, you “might” trip on them,
 - Stalagmite





Cave Formations

- Stalactites
 - Forms for deposits starting on the ceiling
 - They “hold on tight” to the ceiling
 - stalactite





Pre-increment $*++p$

- Pre-increment / pre-decrement
 - Think of it **prematurely** modifying the address.
- It's a bit of a stretch
 - But you'll remember it.





Post-increment `*p++`

`*p++`

- Translates to:

1. `*p`
2. `p = p + 1;`

`*p--`

- Translates to:

1. `*p;`
2. `p = p - 1;`

- Parenthesis does nothing
 - `*(p++) == *p++`
- Remember both
 - post-increment
 - post-decrement
 - Dereferencing happen **before** address processing.
- Faster than pre-increment, (yes!)
 - since the dereference is immediate



Post-increment `*p++`

- Post-increment / post-decrement
 - Give me the pointer NOW
 - We fix it in **POST**
 - modifying the address
- It's a bit of a stretch
 - But you'll remember it.





Big Endian

- Most significant byte is at the first address.
- PowerPC, Motorola Chips, and other processors
- It's the layout I think instantly
 - Makes sense

Big Endian:

unsigned int x = 0xAABBCCDD;

0x00:	AA	BB	CC	DD
-------	----	----	----	----

0x00:	AA
0x01:	BB
0x02:	CC
0x03:	DD



Little Endian

- **Least** significant byte is at the first address.
- Intel, Mips, and other processors
- Backwards, feels weird.
- Created out of legacy issues from the 8088.

Little Endian:

unsigned int x = 0xAABBCCDD;

0x00:	DD	CC	BB	AA
-------	----	----	----	----

0x00:	DD
0x01:	CC
0x02:	BB
0x03:	AA

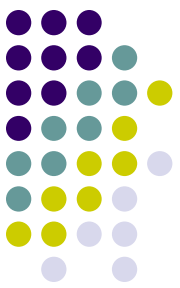


Adding 7 bytes to int *

- Example:
 - `int *p;`
 - `p = p + 7;`
 - Adds 28 bytes to address p not 7
 - Why?
- Don't ask why, here's how you do it.

`p = (int *) ((char *)p + 7);`
or

`p = (int *) ((unsigned int)p + 7);`



Subtracting pointers

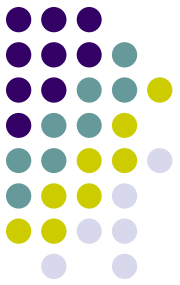
- Can subtract two pointers of the same type.
 - $(p2 - p1)$
- The result is the distance (in array elements) between the two elements.
 - **NOT in BYTES**
- Can be confusing... I convert to `char *` or unsigned ints.
 - `difference_in_bytes = (char *)p2 - (char *)p1;`



Useful information

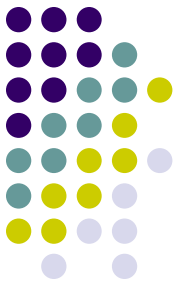
- `printf (“ %p “, p);`
 - `%p` – prints pointers
- Do Not use NULL
 - Use **0 for Null pointers**
 - It's the ANSI spec
- Always test your pointers before dereferencing them

Examples




- On Board:

Thank You!



- Questions?





Simple Memory System

Optimized C++

Ed Keenan



Goals

- Understand the internals of a simple memory system
- Create your own Heap Memory system
 - Not so great party material, too long to describe
 - 3-monkey joke is better



About time!



3 Monkey Joke

- There are 3 monkeys playing in a tree.
 - They climb down from the tree and play in the dirt so they are dirty. **DIRTY MONKEYS**
 - Now the dirty monkeys need to take a bath in the bath tub next to the tree.
- The monkeys are now in the bath tub.
 - The 1st monkey says to the 2nd monkey,
 - **"Pass the brush"**
 - The 2nd monkey says to the 3rd monkey,
 - **"Pass the soap"**
 - And the 1st monkey says to the 3rd monkey,
 - **"Pass the shampoo"**
 - And then the 2nd monkey says,
 - **"What do you think I am, a typewriter?"**




Heaps

- Heap is a stand-alone buffer that contains dynamic allocations
 - Every allocation comes out of one bucket
 - User has control of Heap creation and lifetime
 - Can create memory allocations schemes to suite the code use patterns.
 - It's faster, faster, faster
- This class is all about memory usage is some way or another.
- Macro Benefits
 - Benefits, all the allocations are contained in the bucket.
 - We are not fragmenting other memory buckets.
- Micro Benefits
 - Create different allocation schemes
 - Can be very fast
 - Can have very low fragmentation
 - Can use temporal relationship data
 - Generally faster than global new and delete




Problem of management

- Start off with a buffer, we'll call that the heap.
- Now we need to add usable blocks to the heap.
- How do we manage them?





How do we Free Blocks?




- If we free adjacent blocks
- Coalesce the blocks into one larger free block
 - Gives us more options for future allocations

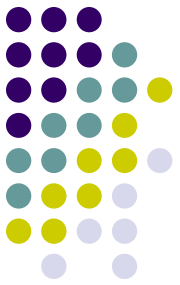


Problem of Fragmentation

- There's empty space, but I can't find a contiguous block available
- Too many allocations sub-dividing the space.



Where's Yojimbo?



- Thanks, now I have a name.





Time to Allocate / Deallocate



- How efficient is the request for allocation, (malloc).
 - Searching for blocks
 - Best strategy
- How much time does it take to free an block?
 - Time is money!
 - We are about speed



What is a Memory System

- Dispenses dynamic memory blocks (*malloc*) and reuses memory blocks discarded through (*free*).
- Physical data structures, management and data layout are very important
- Solving several problems at once:
 - Fragmentation
 - Coalescing of free blocks
 - Searching
 - Algorithms
 - Allocation strategies
 - Management of data structures



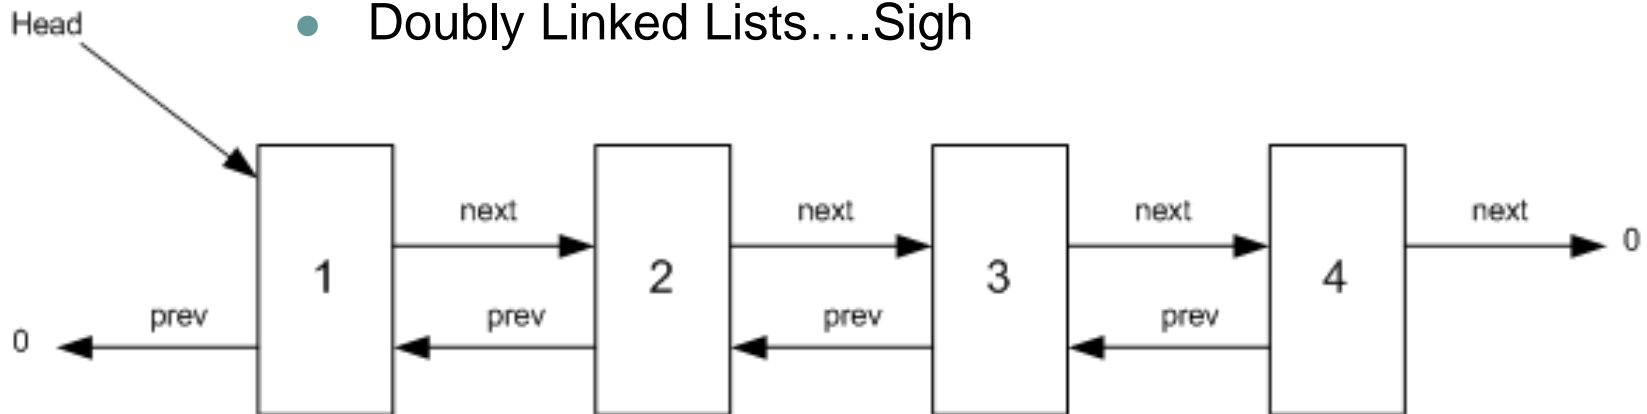
Let's start

- Using *malloc* and *free*
- Yes, I'm old,
 - But that's how it really works at an implementation level
- Operator *new* and *delete* deal with the instantiation of the objects,
 1. Call system malloc / free under the hood.
 2. Call constructors (new, reverse order for delete)



Tracking blocks


- Need to determine which blocks are allocated
- Requirements
 - Easily insert
 - Easy to remove
 - Tracking
- Simple data structure that we all understand
 - Doubly Linked Lists....Sigh




Copernicus




- Does everyone understand doubly linked lists?
 - Speak up.
 - I have a marker board
- Who knows about Copernicus?
 - He's Dead
 - You will be if you don't know these lists inside & out.
- Final Call?






Concept of Headers



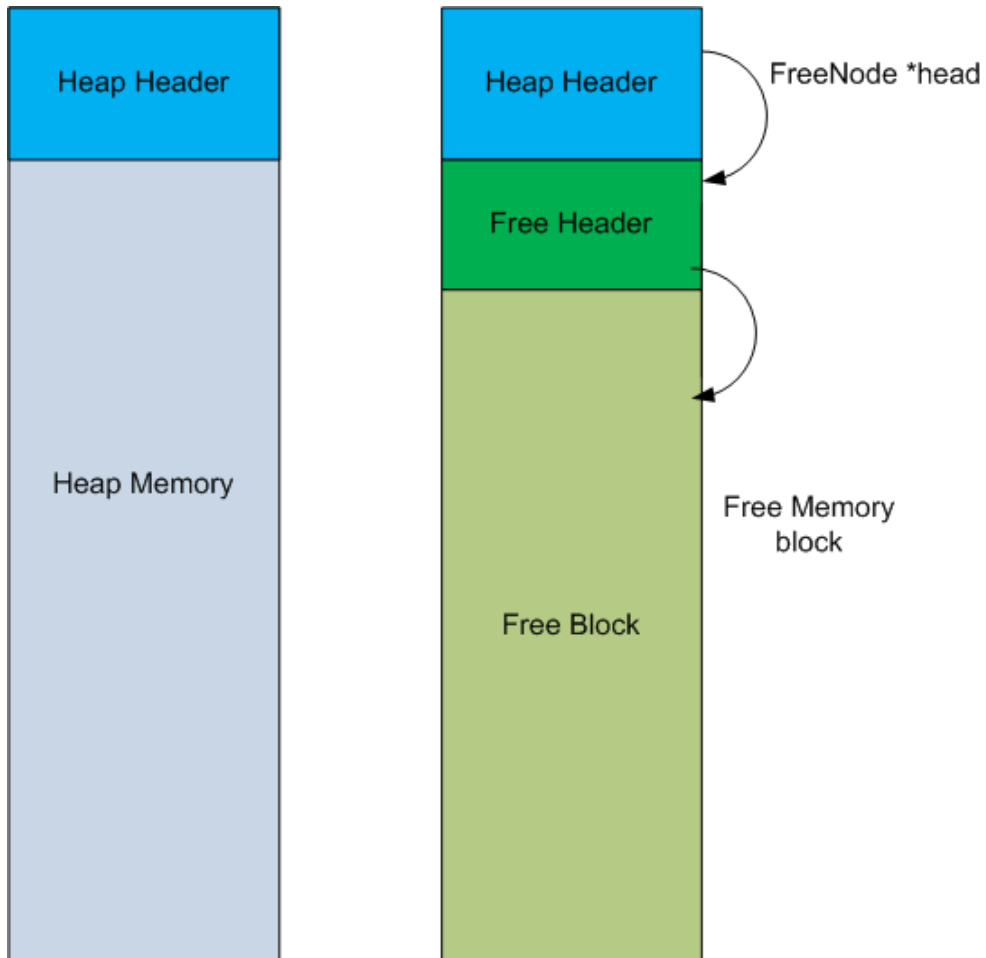
- Disjointed memory
 - Controlling node, dynamic memory allocated elsewhere
 - Remember Hot / Cold structures?
 - McDLT (video please)

- Jointed memory
 - Gluing them together
 - Controlling structure + dynamic memory
- Call Node header





Heap Header

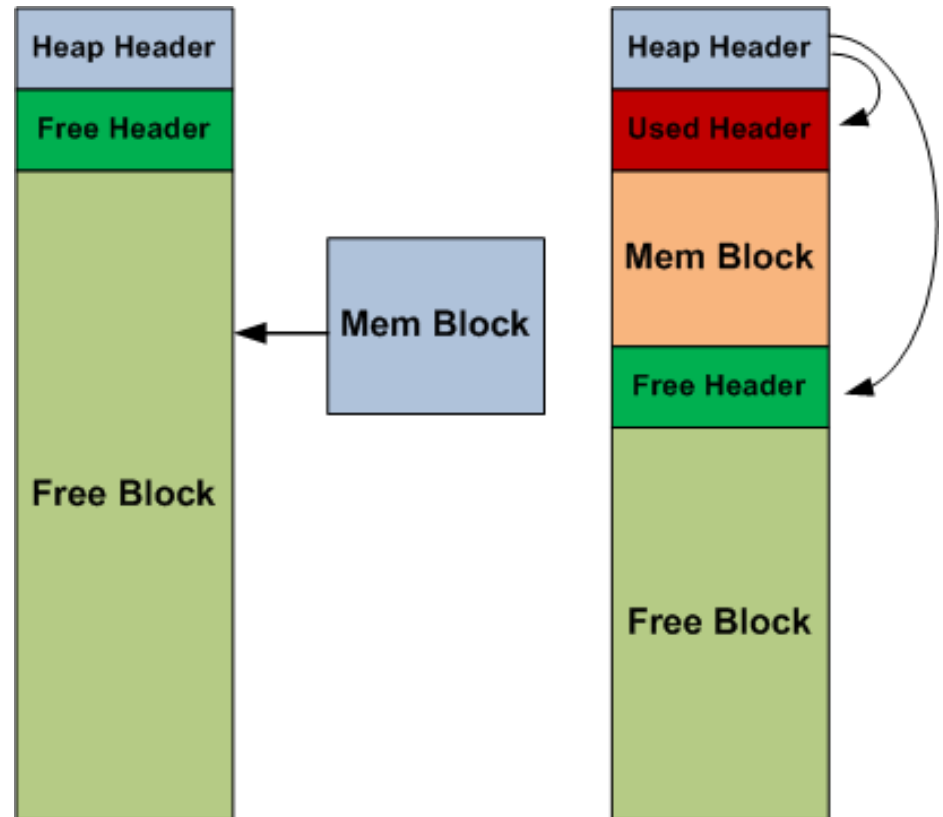


- Heap Header stores
 - Link lists head pointers
 - Tracking data
 - Number allocations
 - Number free blocks
 - Debugging data



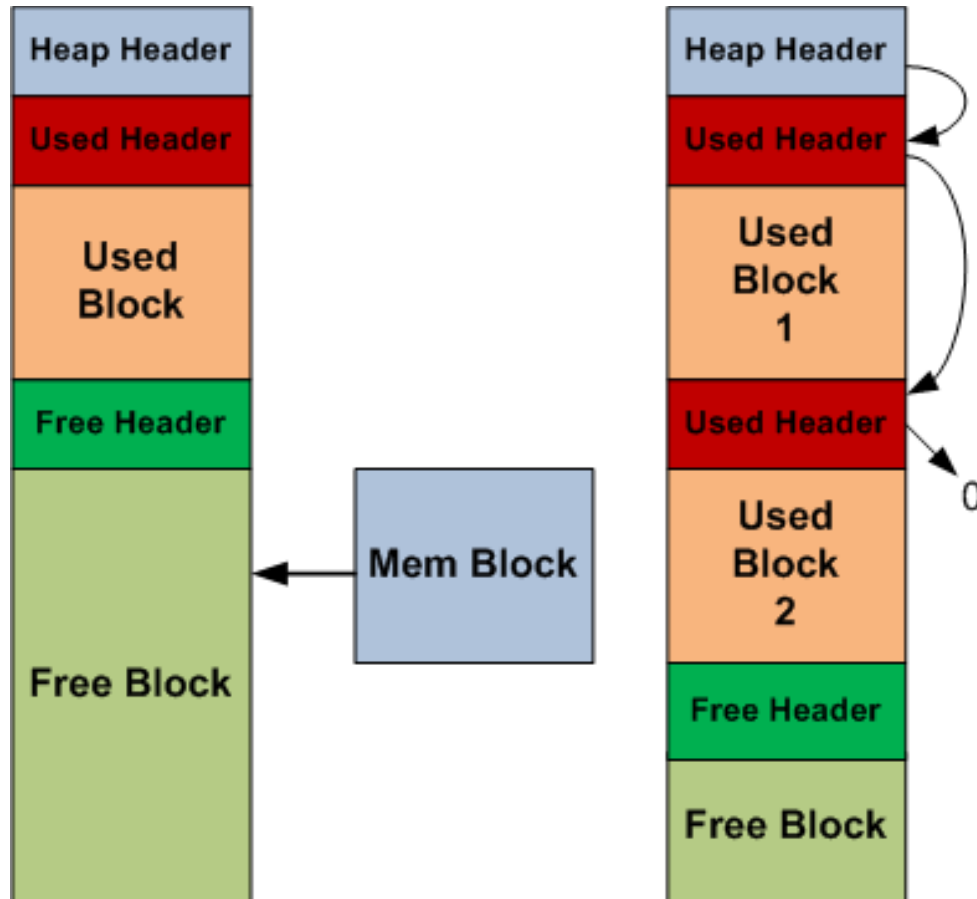
Allocate a block

- Allocating a block from the free block
 - Will the requested block fit?
- Sub-divide the free block
 - Used block & header
 - Free block & header
- Update all the pointers in the heap header





Add another block

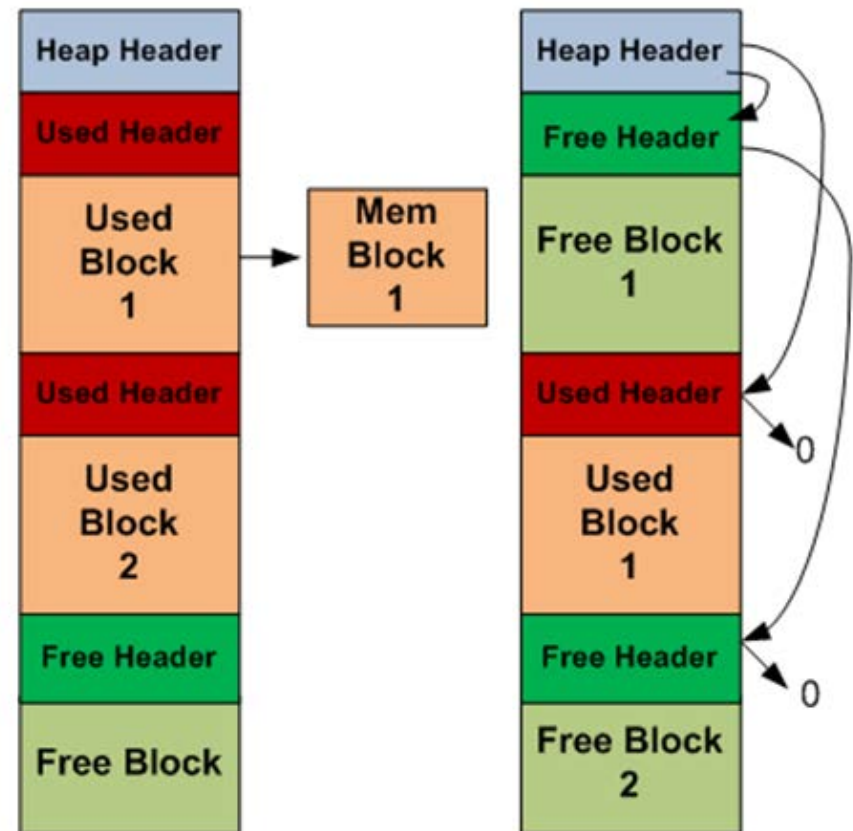


- Add another block
 - Subdivide the free block.
 - Add the 2nd block on to used list.
 - Correct all the linked lists and heap data

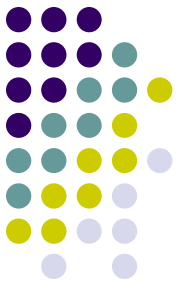


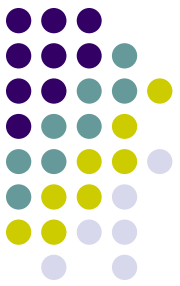
Free a used block

- Free a block
 - Mark the used back as free.
 - Convert the block to free
 - Fix up linked lists

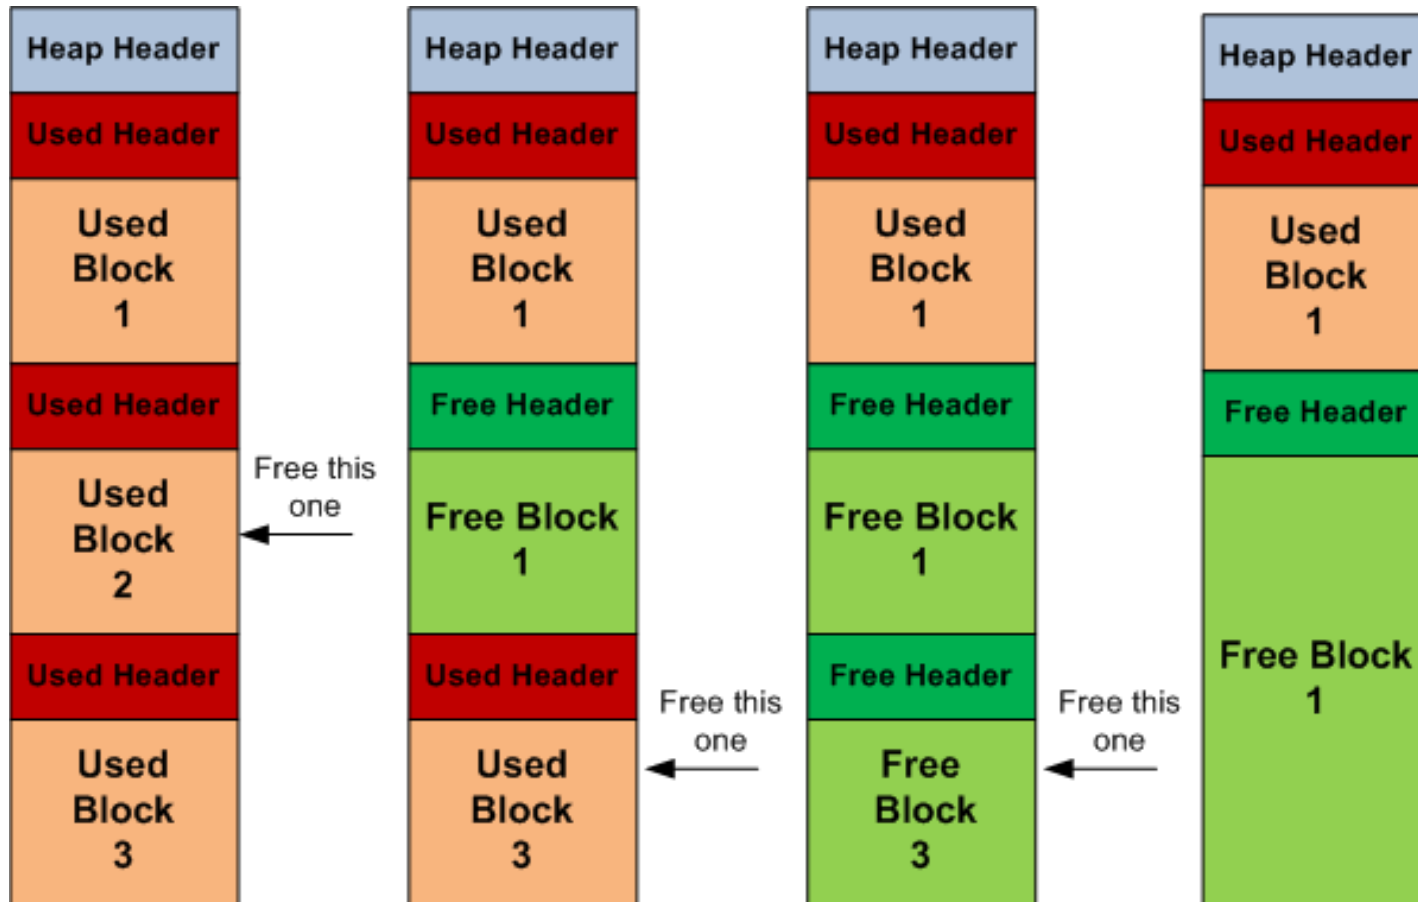


Questions so far






Coalescing



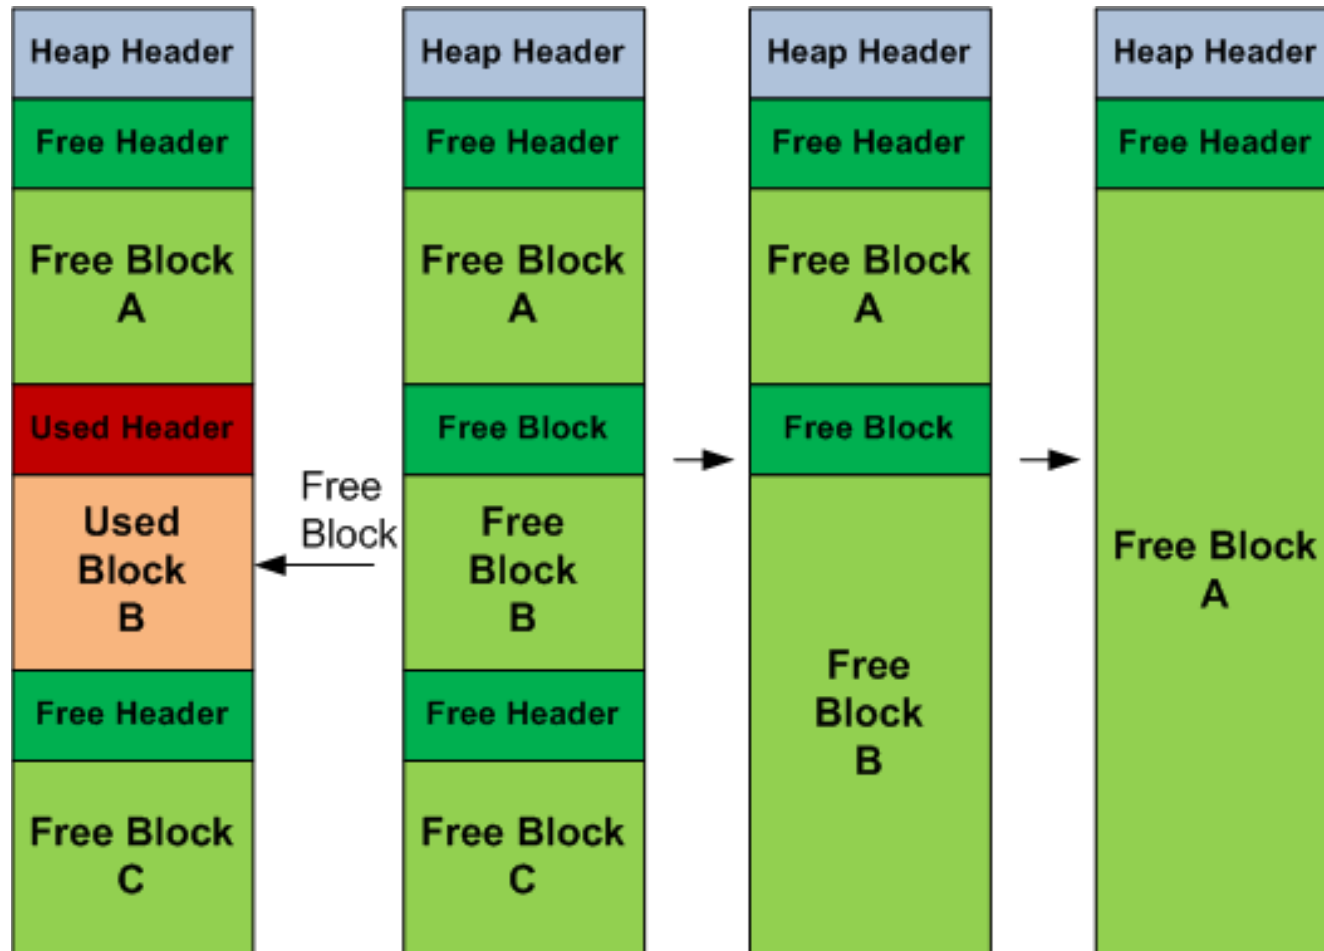


Don't be a cowboy

- Free blocks up in discrete steps
 1. Free the block
 2. Coalesce the blocks
 - Do them one at a time
 - Trust me.
- Kyle's 2 Step Process
 - 6 x better than the famous 12 step process



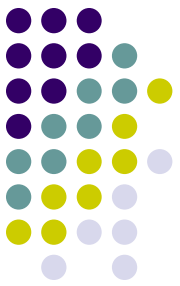
Multiple-sided Coalescing



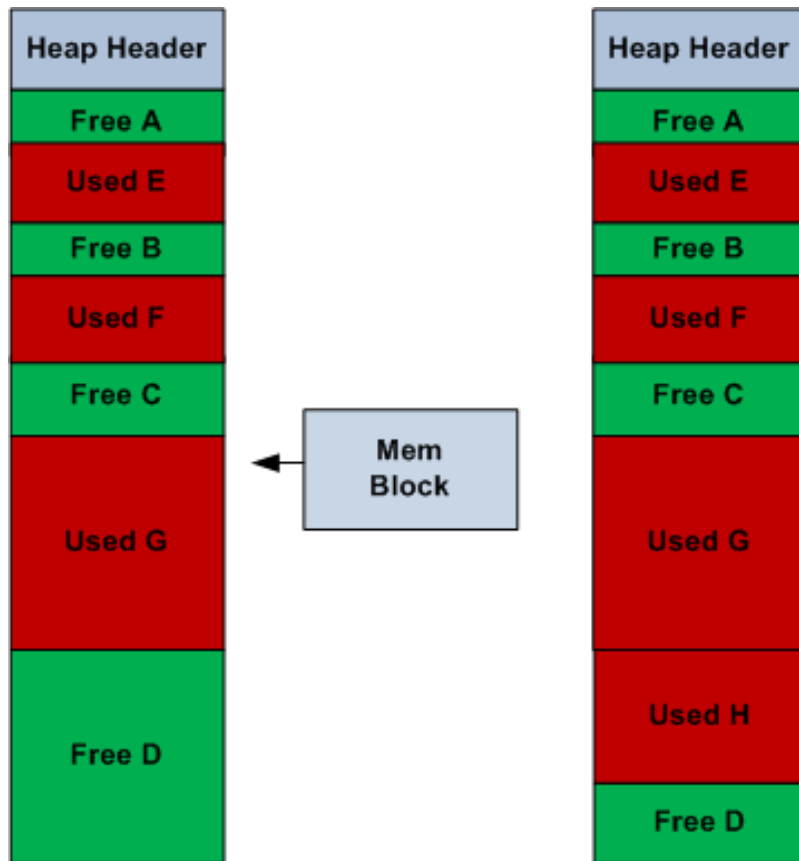


Final thoughts on Coalescing

- Edge conditions will trip you up.
 - My tests will drill into these areas
 - So do the testing.
- Ideas
 - Treat the heap header as a Used Block
 - Some people put another Used Block on the bottom of the heap to reduce edge conditions
 - WE DON'T
- Some books call this compaction



Finding a block



- You might have to search the **Free** linked list until you find a large enough empty block.
- The amount you search directly affects performance.
- Reduce fragmentation
 - Improve performance



Allocation strategies

- Best fit
 - Find the hole the fits the block the best
 - Least amount of waste.
 - Slower
 - Has low fragmentation
- Worse fit
 - Find the biggest **Free** block
 - Take your allocation from that block



Allocation strategies

- First Fit
 - Find the first open block
 - subdivide that block for your allocation
 - Simplest, pretty effective (Use this in class)
- Next Fit
 - Find the 1st hole the fits the block
 - Save the location
 - Start from last location down the heap
 - Wrap back to top
 - Faster
 - → Use this in class PA



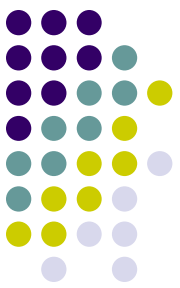
Practical issues

- High level was discussed
- Concepts were conveyed
 - ***Devil lies in the details***
 - Many problems come out of this assignment
- Advice:
 - You need to debug
 - Build and develop incrementally
 - Trust me

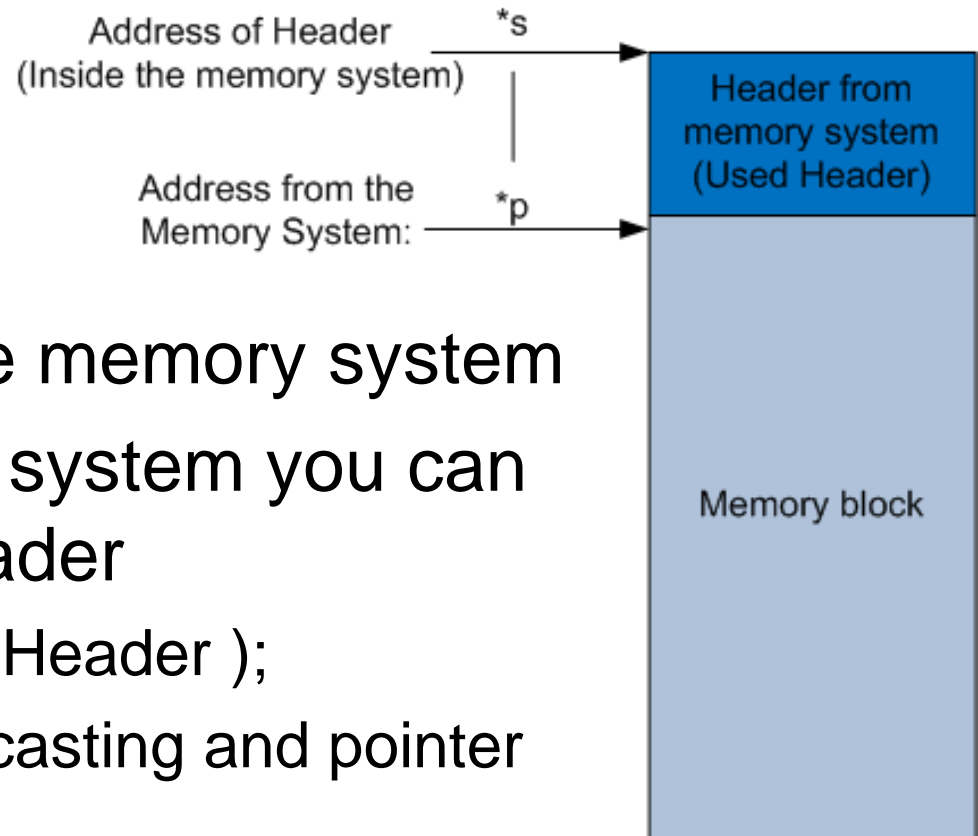


Find Header from Malloc block

- You have a dynamically allocated block
 - Now find the header in the heap
- Naïve way:
 - Search the memory range until you can find the address of the block in heap.
 - Walk the Used pointers until you find address
- Cool way:
 - Move back up from the pointer the size of header.
 - Cast that address to be header
 - Pointer math is useful



Getting Header pointer



- You get $*p$ from the memory system
- Inside the memory system you can get $*s$ the used header
 - $s = p - \text{sizeof}(\text{usedHeader})$;
 - Do forget all those casting and pointer stuff.



Ideas for your headers

- Free Header
 - Node Linkages
 - Next pointer
 - Previous pointer
 - Size
 - Size of free block
 - Flags
 - Miscellaneous flags and fields
 - Any additional data that helps you with your design
- Used Header
 - Node Linkages
 - Next pointer
 - Previous pointer
 - Size
 - Size of free block
 - Flags
 - Flags on adjacency
 - Above block is Free or Used
 - Below block is free or used
 - Miscellaneous flags and fields that help you with your design

Thank You!



- Questions?

